

# 并行与分布式计算

Parallel & Distributed Computing


陈鹏飞  
计算机学院



**中山大学**  
SUN YAT-SEN UNIVERSITY

**计算机学院 (软件学院)**

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



# **Lecture 6 — Race Conditions and Synchronization in OpenMP**

**Pengfei Chen**

**School of Data and Computer Science**




中山大學

SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



## Outline:

- **Correctness issues in parallel programming (in OpenMP)**
  - ❑ **Barriers (障碍, 屏障)**
  - ❑ **Examples of race conditions**
  - ❑ **Mutual Exclusion (互斥)**
  - ❑ **Memory fence (内存屏障)**



# ***Concept: Synchronization***

## ➤ **Synchronization**

- ❑ **The process of managing shared resources so that reads and writes occur in the correct order regardless of how the threads are scheduled**

## ➤ **Synchronization methods**

- ❑ **Barriers**

- ❑ **Mutual Exclusion (互斥) e.g. `pthread_mutex_lock`**

- ❑ **...**



# Barriers in OpenMP

## ➤ Barrier

- ❑ A synchronization point at which every member in a team of threads must arrive before any member can proceed

## ➤ Syntax

#pragma omp barrier

- ❑ Automatically inserted at the end of worksharing constructs
- ❑ e.g., *for* pragma, *single* pragma, ...
- ❑ Can be disabled by using the *nowait* clause



## Example: Use of Barrier

```
int numt = omp_get_num_threads();  
#pragma omp parallel shared(numt)  
{  
    int tid = omp_get_thread_num();  
    printf("hi, from %d\n", tid);  
    if (tid == 0) {  
        printf("%d threads say hi!\n", numt);  
    }  
}
```

Output using 4 threads

hi, from 3

hi, from 0

hi, from 2

4 threads say hi!

hi, from 1

### ◆ Question:

- What's the expected output?
- How can we let the last printf appear last?

Barbara Chapman, "A Guide to OpenMP," 2010.





## *Example: an Explicit Barrier*

```
int numt = omp_get_num_threads();
#pragma omp parallel shared(numt)
{
    int tid = omp_get_thread_num();
    printf("hi, from %d\n", tid);
#pragma omp barrier
    if (tid == 0) {
        printf("%d threads say hi!\n", numt);
    }
}
```

Output using 4 threads

hi, from 3

hi, from 0

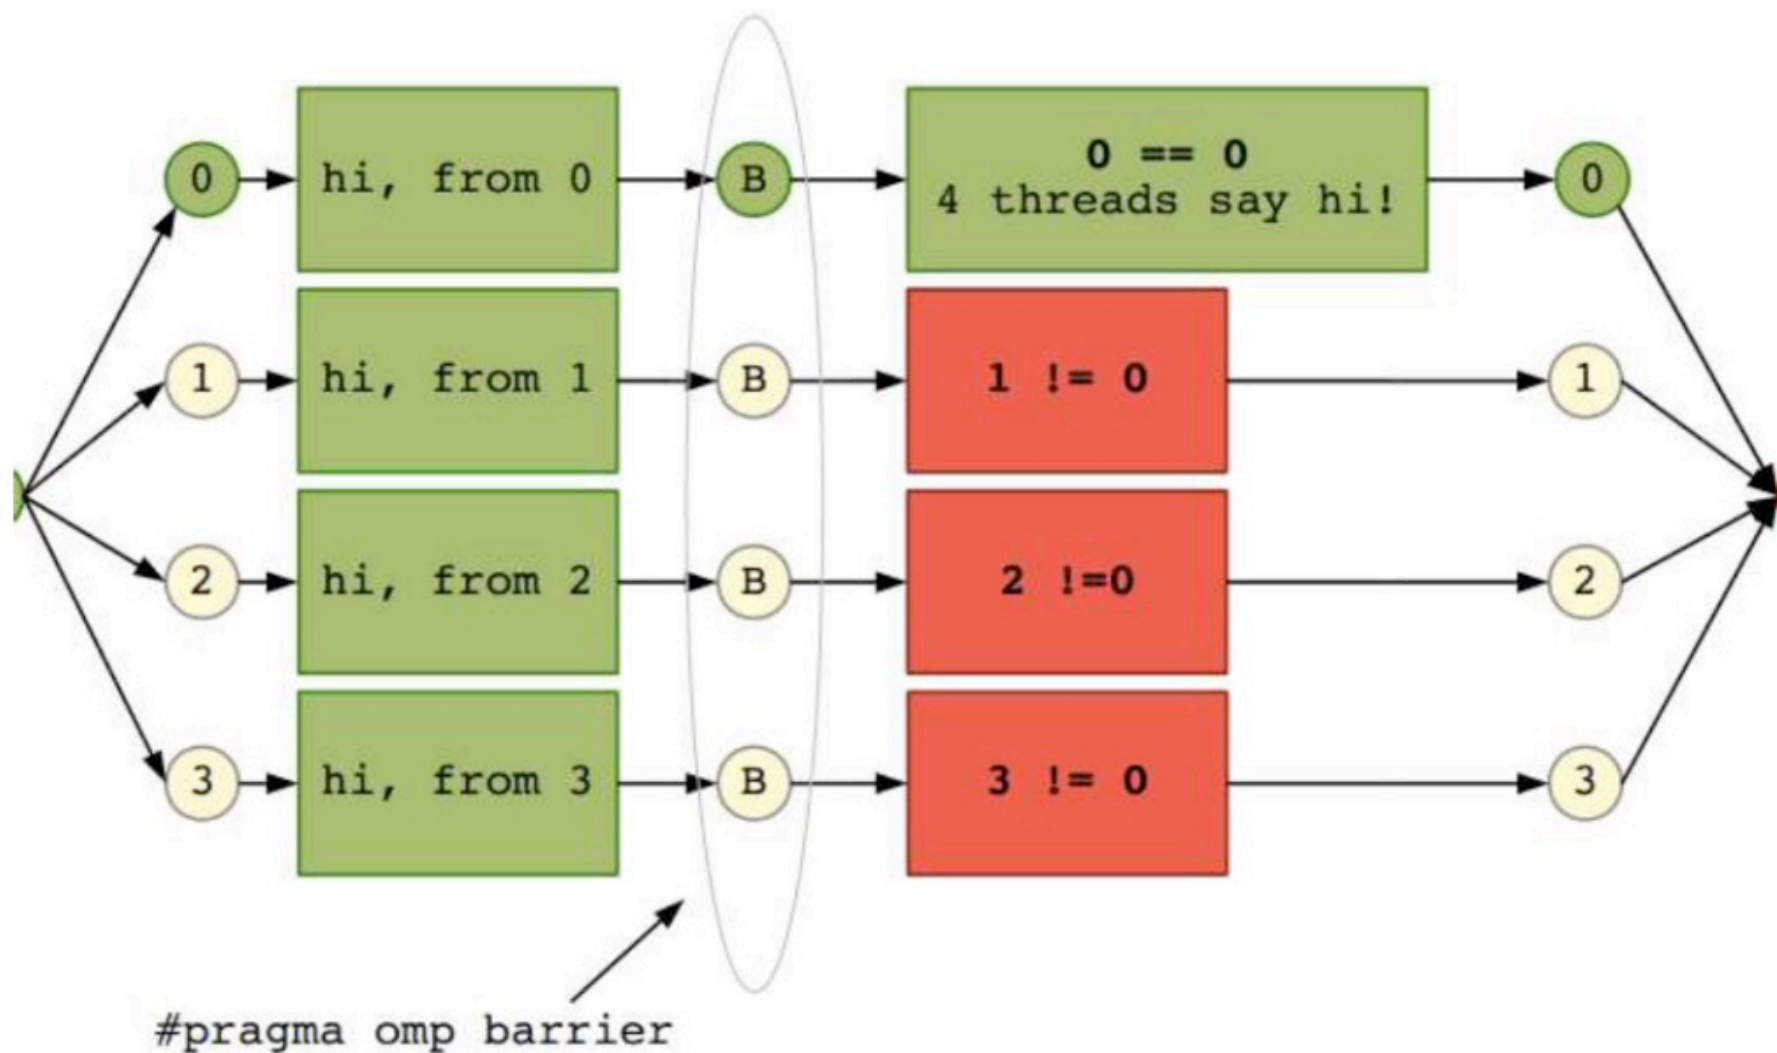
hi, from 2

hi, from 1

4 threads say hi!



## Example: an Explicit Barrier







## ***Clause: nowait***

- The *nowait* clause tells the compiler that there is no need for a barrier

synchronization at the end of a *parallel for* loop or *single* block of code

## ***Case: parallel, for, single Pragmas***

```
for (i = 0; i < N; i++)  
    a[i] = alpha(i);  
if (delta < 0.0)  
    printf("delta < 0.0\n");  
for (i = 0; i < N; i++)  
    b[i] = beta(i, delta);
```



## ***Solution: parallel, for, single Pragmas***

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i = 0; i < N; i++)
        a[i] = alpha(i);
    #pragma omp single nowait
    if (delta < 0.0)
        printf("delta < 0.0\n");
    #pragma omp for
    for (i = 0; i < N; i++)
        b[i] = beta(i, delta);
}
```



# ***Mutual Exclusion***

## ➤ **Mutual exclusion**

- ❑ **A kind of synchronization**

- ❑ **Allows only a single thread or process at a time to have access to shared resource**

- ❑ **Implemented using some form of locking**

## ➤ **Critical section (a high-level synchronization)**

- ❑ **Only one thread at a time will execute the structured block within a critical section**

## ➤ **Lock (a low-level synchronization)**



# *An Example of Race Condition*

```
double area, pni, x;
int i, n;

...

area = 0.0;
for (i = 0; i < n; i++) {
    x = (i + 0.5) / n;
    area += 4.0 / (1.0 + x * x);
}

pi = area / n;
```

➤ What happens when we make the **for** loop parallel?




# Race Condition

- A *race condition* is nondeterministic (非确定性) behavior caused by the times at which two or more threads access a shared variable
- For example, suppose both Thread A and Thread B are executing the statement...

```
area += 4.0 / (1.0 + x*x);
```




## *One Timing $\Rightarrow$ Correct Sum*







## *Another Timing $\Rightarrow$ Incorrect Sum*






## *Another Race Condition Example*

```
struct Node {  
    struct Node* next;  
    int data;  
}  
  
struct List {  
    struct Node* head;  
};  
  
void AddHead (struct List* list, struct Node* node) {  
    node->next = list->head;  
    list->head = node;  
}
```




# *Original Singly-Linked List*






## *Thread 1 after Stmt. 1 of AddHead*






## *Thread 2 Executes AddHead*






## *Thread 1 after Stmt. 2 of AddHead*







# ***Why Race Conditions Are Nasty (令人讨厌)***

- **Programs with race conditions exhibit nondeterministic behavior**
  - ❑ Sometimes give correct result
  - ❑ Sometimes give erroneous result
- **Programs often work correctly on trivial data sets and small number of threads**
- **Errors more likely to occur when number of threads and/or execution time increases**
- **Hence debugging race conditions can be difficult**



# ***How to Avoid Race Conditions***

- **Scope variables to be private to threads**
  - ▣ Use OpenMP *private* clause
  - ▣ Variables declared within threaded functions
  - ▣ Allocate on thread's stack (pass as parameter)
- **Control shared access with critical region**
  - ▣ Mutual exclusion and synchronization



# ***Mutual Exclusion***

- **We can prevent the race conditions described earlier ...**
  - ❑ **Ensure that only one thread at a time references or updates shared variables**
- **Mutual exclusion**
  - ❑ **A kind of synchronization**
  - ❑ **Allows only a single thread or process at a time to have access to shared resource**
  - ❑ **Implemented using some form of locking**



## ***Do Flags Guarantee Mutual Exclusion?***

```
int flag = 0;
```

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```



# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag  
0

Thread 1

```
void AddHead(struct List* list, struct Node* node) {
```

```
    while (flag != 0) /* wait */;
```

```
    flag = 1;
```

```
    node->next = list->head;
```

```
    list->head = node;
```

```
    flag = 0;
```

```
}
```



# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag  
0

Thread 1

Thread 2

```
void AddHead(struct List* list, struct Node* node) {  
  while (flag != 0) /* wait */;  
  flag = 1;  
  node->next = list->head;  
  list->head = node;  
  flag = 0;  
}
```







# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag  
1

```
void AddHead(struct List* list, struct Node* node) {
```

```
    while (flag != 0) /* wait */;
```

```
    flag = 1;
```

```
    node->next = list->head;
```

```
    list->head = node;
```

```
    flag = 0;
```

```
}
```

Thread 1

Thread 2





# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag  
1

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```

Thread 1

Thread 2





# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag

0

Thread 1

Thread 2

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```



# Flags Don't Guarantee Mutual Exclusion

```
int flag = 0;
```

flag  
0

```
void AddHead(struct List* list, struct Node* node) {  
    while (flag != 0) /* wait */;  
    flag = 1;  
    node->next = list->head;  
    list->head = node;  
    flag = 0;  
}
```

Thread 1

Thread 2





# Locking Mechanism

- The previous method fails because...
  - ❑ (i) Checking the value of *flag* and (ii) setting its value are two distinct operations
- We need some sort of *atomic* test-and-set
  - ❑ Operating systems provide functions to do this
- Lock
  - ❑ Synchronization mechanism used to control access to shared resources
  - ❑ (A generic term)



# *Pragma: critical*

- **Critical section**
  - ❑ A portion of code that only one thread at a time may execute  
(mutually exclusive)
- **Syntax in OpenMP**
  - `#pragma omp critical`
- **Good news! (^\_^)**
  - ❑ Critical sections eliminate race conditions
- **Bad news! (@\_@)**
  - ❑ Critical sections are executed sequentially
- **More bad news! (@\_@)**
  - ❑ You have to identify critical sections yourself





## *Is the AddHead() Function Correct Now?*

```
void AddHead(struct List* list, struct Node* node) {  
    node->next = list->head;  
    #pragma omp critical  
    list->head = node;  
}
```



中山大學

SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING




## ***Is the AddHead() Function Correct Now?***

- **You must protect both read and write access to any shared data**
- **For the AddHead() function, both lines need to be protected**



## *Corrected AddHead() Function*

```
void AddHead(struct List* list, struct Node* node) {  
    #pragma omp critical  
    {  
        node->next = list->head;  
        list->head = node;  
    }  
}
```



# ***OpenMP atomic (原子) Construct***

- **Special case of a critical section to ensure atomic update to memory**

**location**

- **Applies only to simple operations:**

- ❑ **Pre- or post-increment (++)**

- ❑ **Pre- or post-decrement (--)**

- ❑ **Assignment with binary operator (of scalar types)**

- **Works on a single statement**

`#pragma omp atomic`

`counter += 5;`



## Critical vs. Atomic

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp critical
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```


### ◆ Critical protect:

- Call to WorkOne()
  - Finding value of index[i]
  - Addition of x[index[i]] and results of WorkOne()
  - Assignment to x array element
- ◆ Essentially, updates to elements in the x array are serialized

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp atomic
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

### ◆ Atomic protects:

- Addition and assignment to x array element
- ◆ Non-conflicting updates with be done in parallel
- ◆ Protection needed only if there are two threads where the index[i] values match



```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
        #pragma omp critical
            x[index[i]] += WorkOne(i);
            y[i] += WorkTwo(i);
    }
}
```

**A**

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
        #pragma omp atomic
            x[index[i]] += WorkOne(i);
            y[i] += WorkTwo(i);
    }
}
```

**B**





## Critical vs. Atomic

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp critical
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

### ◆ Critical protect:

- Call to WorkOne()
  - Finding value of index[i]
  - Addition of x[index[i]] and results of WorkOne()
  - Assignment to x array element
- ◆ Essentially, updates to elements in the x array are serialized

```
#pragma omp parallel for
{
    for (i = 0; i < n; i++) {
#pragma omp atomic
        x[index[i]] += WorkOne(i);
        y[i] += WorkTwo(i);
    }
}
```

### ◆ Atomic protects:

- Addition and assignment to x array element
- ◆ Non-conflicting updates with be done in parallel
- ◆ Protection needed only if there are two threads where the index[i] values match





# Summary

## ➤ Synchronization (in OpenMP)

### □ Barrier

- Statement ordering among different threads
- Any statement after the barrier will be executed after the statements before the barrier in every thread

### □ Mutual Exclusion

- Access ordering of shared resources
- A mechanism to avoid race conditions

### □ Memory fence


- Data-related statement ordering in the same threads
- Any data-related statement before the memory fence will be executed before the statements after the memory fence



中山大學  
SUN YAT-SEN UNIVERSITY

计算机学院 (软件学院)

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



# Thank You !